

## ADBMS Unit 4

### ***Data warehousing and Decision Support: Introduction***

Data warehousing and decision support are crucial components in the field of information technology and business intelligence. Let's break down these concepts to better understand their significance and how they work together.

#### Data Warehousing:

##### Definition:

A data warehouse is a centralized repository that stores large volumes of data from different sources within an organization. It is designed for query and analysis rather than transaction processing. In simpler terms, it's a massive, organized storage system for business data.

##### Key Characteristics:

1. **Subject-Oriented:** Data warehouses are organized around specific subjects, such as sales, customer, or product, to support business analysis needs.
2. **Integrated:** They consolidate data from various sources, including operational databases, into a unified and consistent format.
3. **Time-Variant:** Data warehouses maintain historical data, enabling users to analyze trends and changes over time.
4. **Non-volatile:** Once data is stored in a data warehouse, it is rarely deleted or updated. This ensures a stable and consistent environment for analysis.

##### Purpose:

The primary purpose of a data warehouse is to provide a reliable and efficient platform for reporting and data analysis. It supports decision-making processes by offering a consolidated view of an organization's data, making it easier for users to extract meaningful insights.

#### Decision Support:

##### Definition:

Decision support refers to the use of information and analytical tools to assist decision-making within an organization. Decision support systems (DSS) leverage data, models, and analytical techniques to help users make informed and effective decisions.

##### Key Components:

1. **Data:**
  - Raw data from various sources, including the data warehouse.
  - External data relevant to the decision-making process.
2. **Models:**
  - Mathematical or analytical models that help analyze and simulate different scenarios.

### 3. User Interface:

- The presentation layer that allows users to interact with the system and receive insights.

#### Purpose:

Decision support systems aim to enhance decision-making by providing relevant information, analysis tools, and interactive interfaces. They can be used for various purposes, such as strategic planning, forecasting, and performance analysis.

#### Integration of Data Warehousing and Decision Support:

##### 1. Data Accessibility:

- Data warehouses act as the foundational data storage, providing a single source of truth for decision support systems.

##### 2. Analytical Capabilities:

- Decision support systems leverage the organized and integrated data in a data warehouse to perform complex analyses and generate meaningful insights.

##### 3. Historical Analysis:

- The time-variant nature of data warehouses allows decision support systems to conduct historical analyses, enabling organizations to learn from past trends and performances.

##### 4. User Empowerment:

- The combination of a data warehouse and decision support systems empowers users across various levels of an organization to access, analyze, and interpret data for better decision-making.

In summary, data warehousing and decision support work together to create a robust environment for organizations to manage and analyze their data, ultimately supporting informed decision-making processes.

#### ***OLAP: Multidimensional Data Model***

OLAP, which stands for Online Analytical Processing, is a category of software tools that enable users to interactively analyze multidimensional data from multiple perspectives. OLAP is designed for complex queries and data analysis, providing a more intuitive and efficient way to explore and understand data. One of the key models used in OLAP is the Multidimensional Data Model.

#### Multidimensional Data Model:

In the Multidimensional Data Model, data is organized into a structure that resembles a data cube. This cube is composed of dimensions, measures, and hierarchies. Let's break down these components:

##### 1. Dimensions:

- Dimensions are the categorical attributes or perspectives by which you want to analyze your data. Examples include time, geography, product, and customer. Dimensions provide the context for measures.

## 2. Measures:

- Measures are the numeric values or metrics that you want to analyze. These are the quantitative data points of interest, such as sales revenue, quantity sold, or profit.

## 3. Hierarchies:

- Hierarchies represent the organization of data within each dimension. For example, a time dimension might have hierarchies like year, quarter, month, and day. Hierarchies allow users to drill down or roll up to view data at different levels of detail.

### Characteristics of the Multidimensional Data Model:

#### 1. Data Cube:

- The data cube is a visual representation of multidimensional data, where each cell contains a measure at the intersection of various dimension members.

#### 2. Slicing and Dicing:

- Users can "slice" the cube by selecting a specific value from one dimension, "dice" it by selecting values from two or more dimensions, and "pivot" to view the data from a different perspective.

#### 3. Drill Down and Roll Up:

- Users can drill down to view more detailed data or roll up to see aggregated data at higher levels of abstraction. This flexibility allows for in-depth analysis.

#### 4. Fast Query Performance:

- OLAP databases are optimized for fast query performance, enabling users to interactively explore and analyze data without significant delays.

### Types of OLAP:

#### 1. ROLAP (Relational OLAP):

- ROLAP systems store data in relational databases and use SQL for querying. They create a virtual multidimensional cube by aggregating data dynamically from the relational database.

#### 2. MOLAP (Multidimensional OLAP):

- MOLAP systems store data in a proprietary format optimized for multidimensional analysis. Examples include Microsoft Analysis Services and IBM Cognos TM1.

#### 3. HOLAP (Hybrid OLAP):

- HOLAP systems combine elements of both ROLAP and MOLAP. They store some data in a multidimensional cube and other data in relational tables.

## Benefits of Multidimensional Data Model in OLAP:

### 1. Intuitive Analysis:

- The model provides an intuitive way for users to analyze data, making it easier to understand complex relationships.

### 2. Efficient Aggregation:

- Aggregations are precomputed, leading to faster query performance, especially for summary and aggregate functions.

### 3. Flexible Exploration:

- Users can explore data from various dimensions and hierarchies, facilitating a deeper understanding of business metrics.

In summary, the Multidimensional Data Model in OLAP provides a powerful framework for organizing and analyzing data, offering users a flexible and intuitive way to explore information from different perspectives.

## ***Aggregation Queries***

Aggregation queries are SQL queries that involve the use of aggregate functions to perform calculations on a set of values and return a single aggregated result. These functions operate on groups of rows and produce a summary value for each group. Common aggregate functions include `SUM`, `AVG` (average), `COUNT`, `MIN` (minimum), and `MAX` (maximum). Aggregation queries are often used in databases to analyze and summarize data.

Let's look at a few examples of aggregation queries:

### 1. Simple Aggregation:

Problem: Find the total sales amount for all orders.

```
```sql
SELECT SUM(sales_amount) AS total_sales
FROM orders;
```
```

In this example, the `SUM` function is used to add up the values in the "sales\_amount" column for all rows in the "orders" table.

### 2. Grouping and Aggregation:

Problem: Find the total sales amount for each product category.

```
```sql
```

```
SELECT product_category, SUM(sales_amount) AS total_sales
FROM orders
GROUP BY product_category;
'''
```

Here, the `GROUP BY` clause is used to group the results by the "product\_category" column, and the `SUM` function is applied to calculate the total sales amount for each category.

### 3. Aggregation with Filtering:

Problem: Find the average rating for products with a price greater than \$100.

```
```sql
SELECT AVG(rating) AS average_rating
FROM products
WHERE price > 100;
'''
```

The `AVG` function is used to calculate the average rating for products that meet the specified condition using the `WHERE` clause.

### 4. Counting Records:

Problem: Count the number of orders for each customer.

```
```sql
SELECT customer_id, COUNT(order_id) AS order_count
FROM orders
GROUP BY customer_id;
'''
```

The `COUNT` function is used to count the number of orders for each customer by grouping the results based on the "customer\_id" column.

### 5. Using HAVING with Aggregation:

Problem: Find product categories with an average price greater than \$50.

```
```sql
SELECT product_category, AVG(price) AS average_price
FROM products
GROUP BY product_category
HAVING AVG(price) > 50;
'''
```

The `HAVING` clause is used to filter the results based on the result of an aggregate function. In this example, it filters out product categories with an average price less than or equal to \$50.

These examples demonstrate how aggregation queries can be used to summarize and analyze data in a database. They provide valuable insights into the overall characteristics of the data and are essential tools for reporting and decision support.

### ***Window Queries in SQL***

Window functions, also known as window queries, are a category of functions in SQL that perform a calculation across a specified range of rows related to the current row. These functions are applied to a "window" of rows defined by an OVER clause. Window functions are powerful for analytical queries and are often used in scenarios where you need to perform calculations on a subset of data within a larger result set. Here are some common window functions and examples:

#### **1. ROW\_NUMBER():**

The ROW\_NUMBER() function assigns a unique number to each row within a partition of the result set.

```
```sql
SELECT
  customer_id,
  order_date,
  order_amount,
  ROW_NUMBER() OVER (PARTITION BY customer_id ORDER BY order_date) AS row_num
FROM orders;
```
```

This query assigns a row number to each order within each customer's partition based on the order date.

#### **2. RANK() and DENSE\_RANK():**

RANK() and DENSE\_RANK() functions assign a rank to each row within a partition, with RANK() allowing for ties and DENSE\_RANK() not leaving gaps for ties.

```
```sql
SELECT
  product_id,
  price,
  RANK() OVER (ORDER BY price DESC) AS rank,
  DENSE_RANK() OVER (ORDER BY price DESC) AS dense_rank
FROM products;
```
```

These functions rank products based on their price in descending order.

### 3. LAG() and LEAD():

LAG() and LEAD() functions allow you to access data from preceding or following rows within the result set.

```
```sql
SELECT
  order_id,
  order_date,
  order_amount,
  LAG(order_amount) OVER (ORDER BY order_date) AS previous_order_amount
FROM orders;
```
```

This query shows the order amount of the previous order for each order, ordered by order date.

### 4. SUM() and AVG() Over a Window:

You can use aggregate functions like SUM() and AVG() over a window to calculate cumulative or moving averages.

```
```sql
SELECT
  date,
  revenue,
  SUM(revenue) OVER (ORDER BY date) AS cumulative_revenue,
  AVG(revenue) OVER (ORDER BY date) AS moving_average
FROM daily_sales;
```
```

This query calculates cumulative revenue and a moving average over time.

### 5. NTILE():

NTILE() divides the result set into a specified number of roughly equal groups and assigns a group number to each row.

```
```sql
SELECT
  employee_id,
  salary,
  NTILE(4) OVER (ORDER BY salary DESC) AS salary_quartile
FROM employees;
```

...

This query divides employees into salary quartiles based on their salary in descending order.

Window functions provide a powerful way to perform complex analytical tasks within a SQL query, offering flexibility and efficiency in data analysis and reporting.

### ***Implementation Techniques for OLAP***

Online Analytical Processing (OLAP) involves complex queries and analyses on large volumes of data. Various implementation techniques are employed to optimize OLAP systems for performance, scalability, and user responsiveness. Here are some key implementation techniques for OLAP:

#### **1. Multidimensional Cubes:**

- Description: Data is organized into multidimensional cubes with dimensions representing different aspects of the data (e.g., time, geography, product).
- Benefits: Allows for intuitive exploration and analysis of data from different perspectives. Aggregations are precomputed for faster query response.

#### **2. Aggregation:**

- Description: Precompute and store aggregated values at different levels of granularity in the data cube.
- Benefits: Reduces query response time by providing ready-to-use aggregated results for commonly queried combinations of dimensions.

#### **3. Indexing:**

- Description: Use indexes on key columns to speed up data retrieval. Bitmap indexes are commonly used in OLAP systems.
- Benefits: Accelerates query performance by enabling faster data retrieval based on indexed columns.

#### **4. Partitioning:**

- Description: Partition large tables into smaller, more manageable pieces based on certain criteria (e.g., range, list, hash).
- Benefits: Improves query performance by reducing the amount of data that needs to be scanned for a particular query.

#### **5. Materialized Views:**

- Description: Store the results of a query as a table, often periodically refreshed to reflect changes in the underlying data.
- Benefits: Speeds up query performance by allowing the use of precomputed results for certain queries.

#### **6. Parallel Processing:**

- Description: Distribute query processing across multiple processors or nodes in a parallel or distributed environment.
- Benefits: Increases query performance by leveraging parallelism and utilizing the computing power of multiple resources simultaneously.



#### 7. Cube Compression:

- Description: Reduce the storage space required for cubes by applying compression techniques.
- Benefits: Optimizes storage utilization and can lead to faster data retrieval and reduced disk I/O.

#### 8. Advanced Query Optimization:

- Description: Use advanced query optimization techniques such as query rewriting, query caching, and dynamic query optimization.
- Benefits: Enhances query performance by optimizing query execution plans and reusing previously computed results.

#### 9. In-Memory OLAP:

- Description: Load data into memory for faster access and analysis.
- Benefits: Significantly reduces query response time by eliminating the need to fetch data from disk.

#### 10. Incremental Processing:

- Description: Update the OLAP system incrementally as new data becomes available instead of recomputing the entire cube.
- Benefits: Reduces processing time for updates and ensures that the OLAP system remains responsive.

#### 11. Query Caching:

- Description: Cache the results of frequently executed queries for reuse.
- Benefits: Improves query response time by returning cached results for identical or similar queries.

#### 12. Use of Columnar Storage:

- Description: Store data in a columnar format rather than a row-based format.
- Benefits: Enhances query performance by allowing for more efficient compression, as well as enabling better data pruning and columnar operations.

These implementation techniques are often combined to create a well-tuned and efficient OLAP system that meets the performance requirements of complex analytical queries. The specific techniques employed may vary depending on the OLAP architecture and the characteristics of the underlying data.

### ***Data warehousing***

Data warehousing is a process that involves collecting, storing, managing, and organizing large volumes of data from various sources within an organization. The primary goal of a data warehouse is to provide a centralized and unified view of an organization's data, making it easier for decision-makers to extract insights and make informed decisions. Here are key concepts and components associated with data warehousing:

#### Key Concepts:

##### 1. Data Warehouse:

- A data warehouse is a centralized repository that stores data from different sources in a structured format. It is designed for efficient querying and analysis rather than transaction processing.

## 2. Data Mart:

- A data mart is a subset of a data warehouse that is focused on a specific business line, function, or subject area. Data marts are often created to support the needs of a particular department within an organization.

## 3. ETL (Extract, Transform, Load):

- ETL processes involve extracting data from various source systems, transforming it into a consistent format, and loading it into the data warehouse. This process ensures that data is standardized and suitable for analysis.

## 4. Dimensional Modeling:

- Dimensional modeling is a technique used in data warehousing to organize and model data in a way that is easy to understand and query. It involves creating dimensions (categorical attributes) and facts (numeric measures) to represent business processes.

## 5. Star Schema and Snowflake Schema:

- In dimensional modeling, the star schema and snowflake schema are common database schema designs. The star schema has a central fact table connected to dimension tables, while the snowflake schema normalizes dimension tables into a more structured form.

## 6. Metadata:

- Metadata provides information about the data stored in the data warehouse. It includes details about data sources, transformations, and the structure of the data, helping users understand and manage the data effectively.

## Components of Data Warehousing:

### 1. Data Sources:

- Data warehouses integrate data from various sources, including operational databases, external data feeds, and other systems within the organization.

### 2. Data Warehouse Server:

- The data warehouse server is the central storage and processing unit that manages the data warehouse. It stores and organizes the data for efficient retrieval.

### 3. ETL Tools:

- ETL tools automate the extraction, transformation, and loading processes. Popular ETL tools include Informatica, Talend, and Microsoft SQL Server Integration Services (SSIS).

### 4. OLAP (Online Analytical Processing):

- OLAP tools enable users to interactively analyze and explore data stored in the data warehouse. OLAP supports complex queries and multidimensional analysis.

#### 5. Reporting and Analysis Tools:

- Tools like Tableau, Power BI, and SAP BusinessObjects allow users to create reports, dashboards, and perform ad-hoc analysis on the data stored in the data warehouse.

#### 6. Metadata Repository:

- A metadata repository stores information about the data, including data definitions, relationships, and lineage. It helps users understand and manage the data assets within the data warehouse.

#### 7. Security and Access Control:

- Data warehouses implement security measures to control access to sensitive data. Role-based access controls and encryption are commonly used to secure the data.

#### 8. Data Quality and Governance:

- Data quality processes ensure that data in the warehouse is accurate and consistent. Data governance policies help manage and maintain data quality over time.

Data warehousing plays a crucial role in supporting business intelligence, reporting, and decision-making processes by providing a consolidated and reliable view of an organization's data.

### ***Views and Decision Support***

Views in a database and decision support systems (DSS) are closely related, as views provide a mechanism to present a subset of data to users, and decision support systems leverage these views for analysis and decision-making. Let's explore how views are used in decision support:

#### Views in a Database:

##### 1. Definition:

- In a database, a view is a virtual table based on the result of a SELECT query. It does not store the data itself but provides a way to represent a specific subset of data from one or more tables.

##### 2. Purpose:

- Views are created to simplify complex queries, provide a security layer by restricting access to certain columns or rows, and present a specific perspective on the data.

##### 3. Example:

- Suppose you have a database with tables for "Employees" and "Departments." You could create a view that combines information from these tables, showing only the relevant columns and rows, such as employee name, department, and job title.

```
``sql
```

```
CREATE VIEW EmployeeDetails AS
```

```
SELECT EmployeeID, EmployeeName, DepartmentName, JobTitle
```

```
FROM Employees
```

```
JOIN Departments ON Employees.DepartmentID = Departments.DepartmentID;
```

...

## Views in Decision Support:

### 1. Data Abstraction:

- Views in a decision support context provide a layer of abstraction over the raw data, presenting information in a way that is meaningful for analysis. Decision-makers often don't need to interact directly with the underlying tables but can work with pre-defined views tailored to their needs.

### 2. Performance:

- Views can enhance performance by aggregating or summarizing data in advance. Decision support systems often involve complex queries, and views allow for the creation of pre-aggregated tables, making query execution more efficient.

### 3. Security:

- Views contribute to data security by restricting access to sensitive information. Decision support views can be designed to include only the relevant data for analysis, ensuring that users don't access or inadvertently alter data they shouldn't.

### 4. Customization:

- Decision support views can be customized to meet specific reporting and analysis requirements. They can include calculated fields, filters, and transformations that align with the analytical needs of the users.

## Use Cases:

### 1. Executive Dashboards:

- Views can be created to support executive dashboards, providing summarized and aggregated information about key performance indicators (KPIs) without exposing detailed operational data.

### 2. Financial Reporting:

- Decision support views can be designed to cater to financial reporting requirements, presenting financial data in a format that is compliant with accounting standards and regulations.

### 3. Operational Analytics:

- Views can be utilized to support operational analytics by presenting real-time or near-real-time data for monitoring and decision-making in day-to-day operations.

### 4. Ad-Hoc Analysis:

- Decision support views empower users to perform ad-hoc analysis by providing a foundation for exploring data without the need to construct complex queries directly on raw tables.

In summary, views in a database serve as a crucial tool for decision support by providing a tailored, abstracted, and secure way to access and analyze data. They contribute to the effectiveness of decision support systems by simplifying data access and enhancing performance for analytical tasks.

## ***View Materialization***

View materialization is a database optimization technique where the result of a view, which is typically a virtual representation of data based on a query, is stored persistently as a physical table. Instead of computing the view's result on the fly every time it is queried, the precomputed data is stored, improving query performance at the cost of storage and potential staleness of the data.

### Benefits of View Materialization:

#### 1. Improved Query Performance:

- Materializing a view allows for faster query execution because the result set is already computed and stored. This is particularly beneficial for complex queries or views involving aggregations.

#### 2. Reduced Computational Overhead:

- Materialization reduces the need to compute the view's result set on-the-fly, saving computational resources, especially in scenarios where the view is queried frequently.

#### 3. Consistent Results:

- Materialized views provide consistent results over time. Since the data is precomputed and stored, users will always retrieve the same information until the materialized view is refreshed.

#### 4. Enhanced Stability:

- Materialized views can contribute to system stability by reducing the load on the database during peak usage times, as the computed results are stored and readily available.

### Challenges and Considerations:

#### 1. Storage Overhead:

- Materializing views consumes additional storage space. The storage requirements depend on the size of the view and how frequently it needs to be refreshed.

#### 2. Data Staleness:

- Depending on the refresh strategy, materialized views might not reflect real-time data. Users need to be aware of potential delays between updates to the underlying data and the refresh of the materialized view.

#### 3. Maintenance Overhead:

- Materialized views require maintenance, especially when the underlying data changes. Regular updates or refreshes are needed to keep the materialized view synchronized with the source data.

#### 4. Selecting the Right Views:

- Not all views benefit from materialization. Views that are queried frequently and involve complex computations or aggregations are good candidates, but careful consideration is needed to determine which views to materialize.

## Refresh Strategies:

### 1. On-Demand Refresh:

- Refresh the materialized view manually when needed. This gives users control over when the view is updated but may result in slightly outdated data.

### 2. Scheduled Refresh:

- Automate the refresh process on a predefined schedule (e.g., daily, hourly). This ensures regular updates but may lead to more frequent staleness in the data.

### 3. Incremental Refresh:

- Update only the changed or new data since the last refresh. This strategy is useful for large datasets where a full refresh is resource-intensive.

## Example SQL for Materialized View Creation:

```
```sql
CREATE MATERIALIZED VIEW my_materialized_view AS
SELECT column1, column2, aggregate_function(column3)
FROM my_table
GROUP BY column1, column2;
```
```

This SQL statement creates a materialized view by selecting specific columns and applying an aggregate function to one of the columns from the underlying table.

Materialized views are a trade-off between improved query performance and the additional storage and maintenance overhead. The decision to use materialized views should be based on the specific requirements and characteristics of the data and queries in a given database system.

## ***Maintaining Materialized Views***

Maintaining materialized views involves keeping the data in the view synchronized with changes in the underlying base tables. Since materialized views store precomputed results, it's essential to update them to reflect any modifications, additions, or deletions in the source data. Here are key considerations and strategies for maintaining materialized views:

### 1. Full Refresh:

- Description: Recompute the entire materialized view by executing the original query.
- Pros:
  - Simple and guarantees complete synchronization.
  - Suitable for relatively small datasets or infrequently updated views.
- Cons:
  - Resource-intensive for large datasets.

- May result in longer periods of unavailability during the refresh.

## 2. Incremental Refresh:

- Description: Update only the changed or new data since the last refresh.
- Pros:
  - More efficient for large datasets with relatively few changes.
  - Reduces computational overhead compared to a full refresh.
- Cons:
  - Requires tracking changes, which could add complexity.
  - The incremental approach may not be suitable for certain types of views.

## 3. Scheduled Refresh:

- Description: Automate the refresh process on a predefined schedule (e.g., daily, hourly).
- Pros:
  - Regular updates to keep the materialized view relatively current.
  - Can be less resource-intensive than a full refresh.
- Cons:
  - May result in occasional staleness between refresh intervals.
  - Scheduling must align with business needs and system resource availability.

## 4. On-Demand Refresh:

- Description: Refresh the materialized view manually when needed.
- Pros:
  - Users have control over when the view is updated.
  - Can be suitable for less frequently accessed views.
- Cons:
  - Requires user intervention, which may lead to outdated data.
  - May not be ideal for real-time or near-real-time reporting.

## 5. Use of Triggers:

- Description: Use database triggers to automatically refresh the materialized view when certain events occur (e.g., data changes in base tables).
- Pros:
  - Provides an automatic and event-driven way to keep the materialized view updated.
- Cons:
  - Adds complexity to the database schema.
  - Careful management is needed to avoid performance implications.

## 6. Fast Refresh:

- Description: Some databases support fast refresh options, where only the changed data is processed.
- Pros:
  - Faster than a full refresh and more efficient for large datasets.
- Cons:

- Requires specific conditions (e.g., use of materialized view logs) and may not be supported for all types of views.

#### 7. Logging Changes:

- Description: Maintain logs of changes to the base tables to track modifications.
- Pros:
  - Facilitates incremental refresh by identifying changes.
  - Reduces the need for scanning entire tables during refresh.
- Cons:
  - Adds overhead to track and maintain change logs.

#### 8. Maintaining Indexes and Statistics:

- Description: Regularly update indexes and statistics on materialized views to ensure optimal query performance.
- Pros:
  - Improves query performance.
  - Ensures that the database optimizer makes informed decisions.
- Cons:
  - Adds overhead during maintenance operations.

The choice of a maintenance strategy depends on factors such as the size of the dataset, frequency of updates, query performance requirements, and business needs. Combining different strategies based on the characteristics of the materialized view and the underlying data is often a practical approach. Regular monitoring and optimization are crucial for maintaining materialized views efficiently over time.